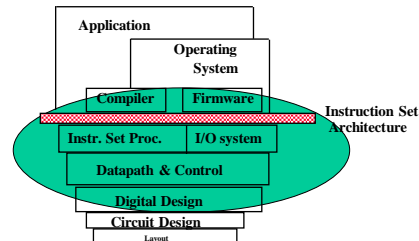
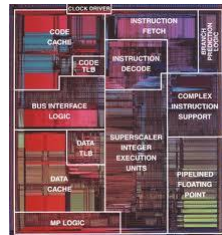


CS/SE 3340

Computer Architecture




Pipelining and Hazards

Adapted from slides by Profs. D. Patterson and J. Hennessey

Questions

1. What/Why of pipelining?
2. How does MIPS design use pipelining ?
3. What/Why of hazards?
4. What are the three types of hazards?
5. How to deal with hazards?
6. What is 'branch prediction'?

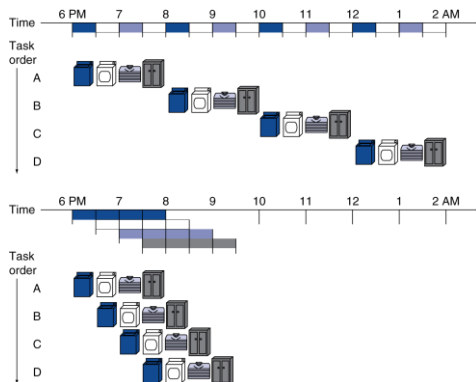


Performance Issues

- Longest delay determines clock period
 - **Critical path**: load instruction
 - *Instruction memory* → *register file* → *ALU* → *data memory* → *register file*
- Not feasible to vary period for different instructions
- Violates design principle
 - Making the common case fast
- We will improve performance by **pipelining**

Pipelining Analogy

- Pipelined laundry: *overlapping execution*
– (temporal) Parallelism improves performance



- Four loads:
 - Speedup
= $8/3.5 = 2.3$
- Non-stop:
 - Speedup
= $2n/0.5n + 1.5 \approx 4$
= number of stages

5

MIPS Pipeline

- Five stages, one step per stage
 1. IF: Instruction fetch from memory
 2. ID: Instruction decode & register read
 3. EX: Execute operation or calculate address
 4. MEM: Access memory operand
 5. WB: Write result back to register

6

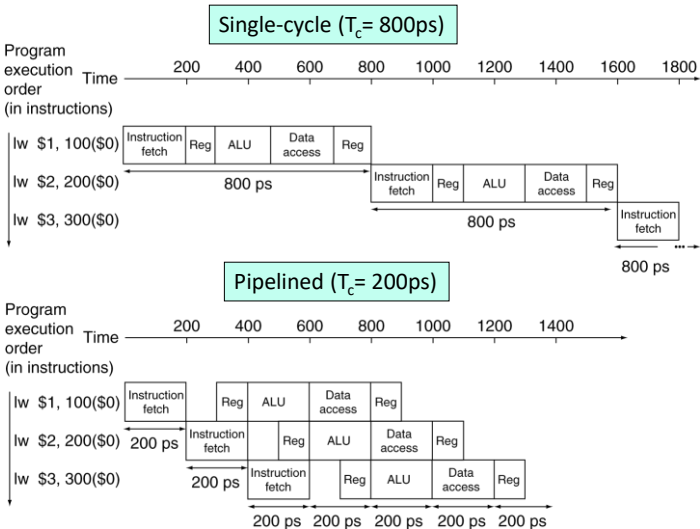
Pipeline Performance

- Assume time for stages is
 - 100ps for register read or write
 - 200ps for other stages
- Compare pipelined datapath with single-cycle datapath

Instr	Instr fetch	Register read	ALU op	Memory access	Register write	Total time
lw	200ps	100 ps	200ps	200ps	100 ps	800ps
sw	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps

7

Pipeline Performance



8

Pipeline Speedup

- If all stages are balanced
 - i.e., all take the same time
 - Time between instructions_{pipelined}
= $\frac{\text{Time between instructions}_{\text{nonpipelined}}}{\text{Number of stages}}$
- If not balanced, speedup is less
- Speedup due to increased throughput
 - Latency (time for each instruction) *does not* decrease

9

Pipelining and ISA Design

- MIPS ISA designed for *pipelining*
 - All instructions are 32-bits
 - Easier to fetch and decode in one cycle
 - c.f. x86: 1- to 17-byte instructions
 - Few and regular instruction formats
 - Can decode and read registers in one step
 - Load/store addressing
 - Can calculate address in 3rd stage, access memory in 4th stage
 - Alignment of memory operands
 - Memory access takes only one cycle

10

Hazards

- Situations that prevent starting the next instruction in the next cycle
- *Structure hazards*
 - A required resource is busy
- *Data hazard*
 - Need to wait for previous instruction to complete its data read/write
- *Control hazard*
 - Deciding on control action depends on previous instruction

11

Structure Hazards

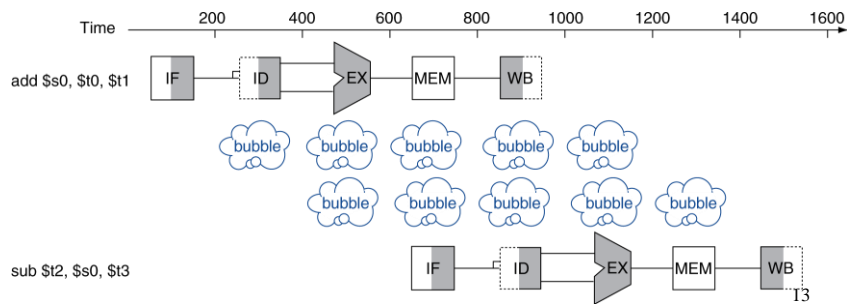
- Conflict for use of a resource
- In MIPS pipeline with a single memory
 - Load/store requires data access
 - Instruction fetch would have to *stall* for that cycle
 - Would cause a pipeline “bubble”
- Hence, pipelined datapaths require separate instruction/data memories
 - Or separate instruction/data caches

12

Data Hazards

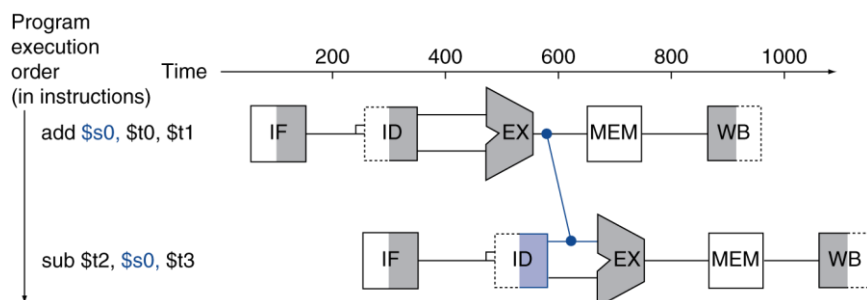
- An instruction depends on completion of data access by a previous instruction

– add \$s0, \$t0, \$t1
 sub \$t2, \$s0, \$t3



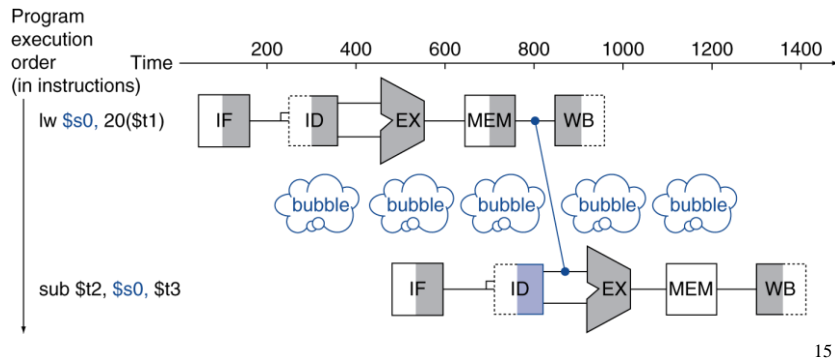
Forwarding (aka Bypassing)

- Use result when it is computed
 - Don't wait for it to be stored in a register
 - Requires extra connections in the datapath



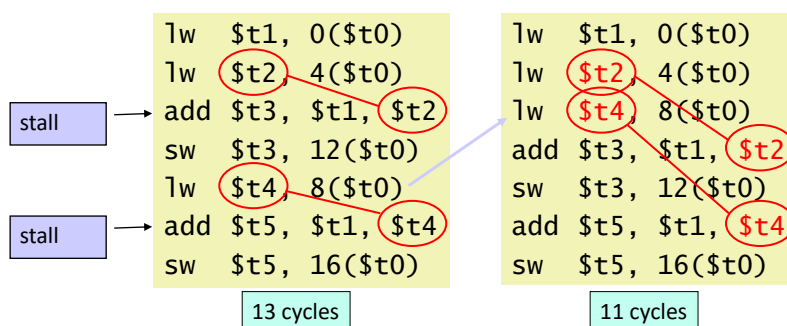
Load-Use Data Hazard

- Can't always avoid stalls by forwarding
 - If value not computed when needed
 - Can't forward backward in time!



Code Scheduling to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction
- C code for $A = B + E$; $C = B + F$;



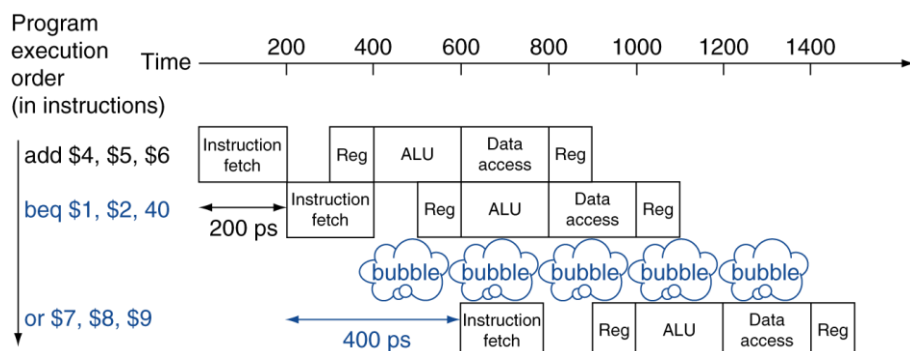
Control Hazards

- Branch determines flow of control
 - Fetching next instruction depends on branch outcome
 - Pipeline can't always fetch correct instruction
 - Still working on ID stage of branch
- In MIPS pipeline
 - Need to compare registers and compute target early in the pipeline
 - Add hardware to do it in ID stage

17

Stall on Branch

- Wait until branch outcome determined before fetching next instruction



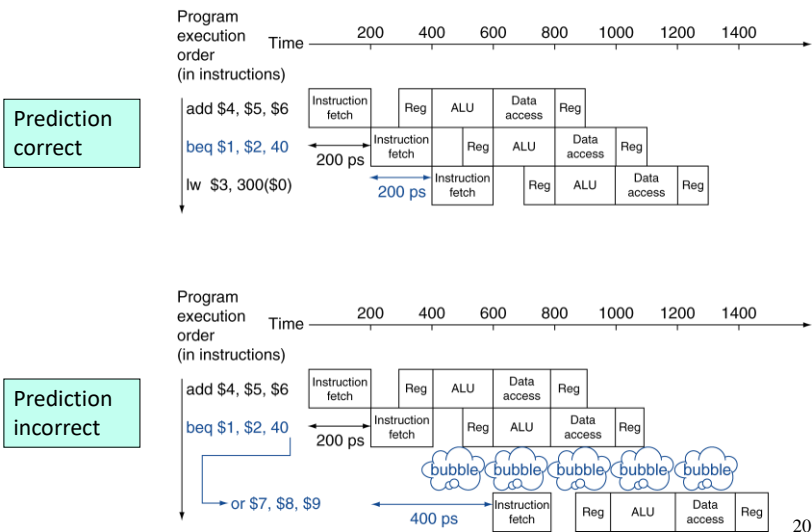
18

Branch Prediction

- Longer pipelines can't readily determine branch outcome early
 - Stall penalty becomes unacceptable
- Predict outcome of branch
 - Only stall if prediction is wrong
- In MIPS pipeline
 - Can predict branches not taken
 - Fetch instruction after branch, with no delay

19

MIPS with Predict: Not Taken



More-Realistic Branch Prediction

- *Static branch prediction*
 - Based on typical branch behavior
 - Example: loop and if-statement branches
 - Predict backward branches taken
 - Predict forward branches not taken
- *Dynamic branch prediction*
 - Hardware measures actual branch behavior
 - e.g., record recent history of each branch
 - Assume future behavior will continue the trend
 - When wrong, stall while re-fetching, and update history

21

Summary

- Modern processors take advantages of the “pipelined effect” to achieve higher performance
 - Instruction execution is divide into stages (ala assembly line)
 - If all stages are balanced at all time, an n-times speed up can be achieved (n: number of stages)
- Hazards cause stages of instruction execution pipeline to be stalled
- Various techniques have been developed to deal with hazards

22